

Document Review: test guide

Upload PDFs 01, 02 and 03 together. Keep this guide separate so the agent has to find the evidence itself. These are development fixtures, not a benchmark.

1. Comparison

Ask: Which model has lower location error and which is faster? Use final results and cite the pages.

Expected: Candidate B has lower RMSE (0.71 m versus 0.84 m). Baseline A has lower mean latency (180 ms versus 260 ms). The reports use the same 90 test cases and workstation; exact hardware is unknown. Evidence: Baseline p2; Candidate p2-p3.

2. Conflicting evidence

Ask: The candidate report says both 0.62 and 0.71 metres. Which value should I use, and why?

Expected: Use 0.71 m. Candidate p3 explicitly replaces the draft subset score of 0.62 m on p1. A good answer acknowledges both values and cites the correction.

3. Release decision

Ask: Does either model meet all release requirements? Identify failed or missing criteria.

Expected: Neither. A fails accuracy ($0.84 > 0.75$) but passes mean latency ($180 \leq 200$). B passes accuracy ($0.71 \leq 0.75$) but fails mean latency ($260 > 200$). Both lack sensor-failure evaluation/recovery and a field pilot. Evidence: Checklist p1; Baseline p2-p3; Candidate p2-p3.

Check abstention and citations

4. Missing information

Ask: Which GPU and how much RAM were used? Was B's improvement statistically significant?

Expected: The exact GPU/RAM and statistical significance are not established. Do not invent hardware specifications, confidence intervals or a p-value.

5. Unsupported generalization

Ask: Can I conclude that Candidate B is robust to hydrogen service and missing sensors?

Expected: No. These conditions were not tested. The agent should distinguish missing evidence from a demonstrated failure.

Inspect the sources. Open each citation and check that the quoted passage supports the claim. A valid page link alone is not enough. Compare the final answer with all three PDFs, including the correction.

Local check. After setup, disconnect external networking and repeat a question. Confirm the app uses local mode. Successful offline operation is a useful deployment check, not a comprehensive security audit.

Generated fictional material for repeatable manual testing. No actual research results or personal information are included.